

Medium

Search



★ Get unlimited access to the best of Medium for less than \$1/week. [Become a member](#)



# Receptive Field



Saba Hesarakhi · Follow


2 min read · Oct 18, 2023

Listen

Share

More

What does the *Receptive field* mean in *CNN* computer vision?



In Convolutional Neural Networks (CNNs) used in computer vision, the receptive field refers to the region of the input image that a particular neuron in a convolutional layer is “looking at” or taking into account when making its predictions or feature extractions. It is a concept used to understand the context and spatial coverage of neurons within the network.

The receptive field is determined by the size of the convolutional kernel (also known as the filter) and the stride used during the convolution operation. The receptive field of a neuron in a particular layer can be thought of as the area of the input image that directly influences the neuron’s output.

Here are some key points to understand about receptive fields in CNNs:

1. **Local Receptive Field:** In the initial layers of a CNN, individual neurons have a small local receptive field, which means they only look at a small portion of the input image. As you move deeper into the network, the receptive field of neurons typically grows, allowing them to capture information from larger regions of the input.
2. **Spatial Hierarchy:** CNNs are designed in a hierarchical manner, where each layer captures features from a larger portion of the input than the previous layer. This hierarchy helps in recognizing both low-level features (e.g., edges and textures) and high-level features (e.g., object parts and objects) as you progress through the layers.
3. **Receptive Field Size Calculation:** The size of the receptive field for a neuron in a given layer depends on the size of the convolutional kernel and the strides used. If a layer uses larger kernels and strides, its neurons have larger receptive fields. Conversely, smaller kernels and strides result in smaller receptive fields.
4. **Padding:** Padding is often used to control the size of the output feature maps. Zero-padding the input can be employed to ensure that the output feature maps have the same spatial dimensions as the input. This can be useful for maintaining the spatial hierarchy and avoiding a reduction in spatial resolution too quickly.

Understanding the receptive field in a CNN is crucial because it helps in interpreting what information is being captured by different layers and neurons in the network. As layers capture increasingly larger receptive fields, they can learn to recognize more complex and abstract patterns, making them suitable for tasks like object recognition and scene understanding in computer vision.


Convolutional Network

Deep Learning

Computer Vision

Kernel

Receptive Field

[Follow](#)

## Written by Saba Hesaraki

175 Followers · 19 Following

MS.c On Artificial Intelligence




No responses yet




What are your thoughts?

[Respond](#)

## More from Saba Hesaraki





 Saba Hesaraki

## Feature Map

What does Feature Map mean in CNN computer vision?

Oct 18, 2023  28

 Saba Hesaraki


## 3D CNN

A 3D Convolutional Neural Network (3D CNN) is a type of deep learning model used for image segmentation in three-dimensional data, such as...

Nov 11, 2023  13



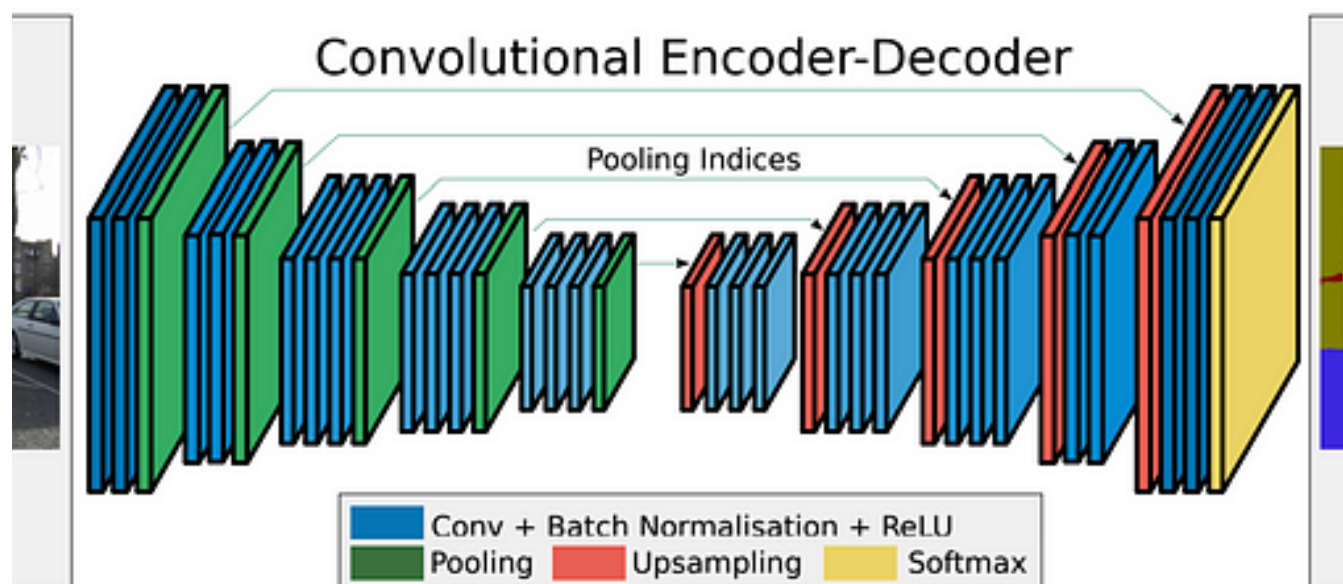



 Saba Hesaraki

## Tversky loss

The Tversky loss is a loss function used in the context of image segmentation tasks, particularly in medical image analysis and computer...

Oct 27, 2023  7  2



 Saba Hesaraki

## SegNET

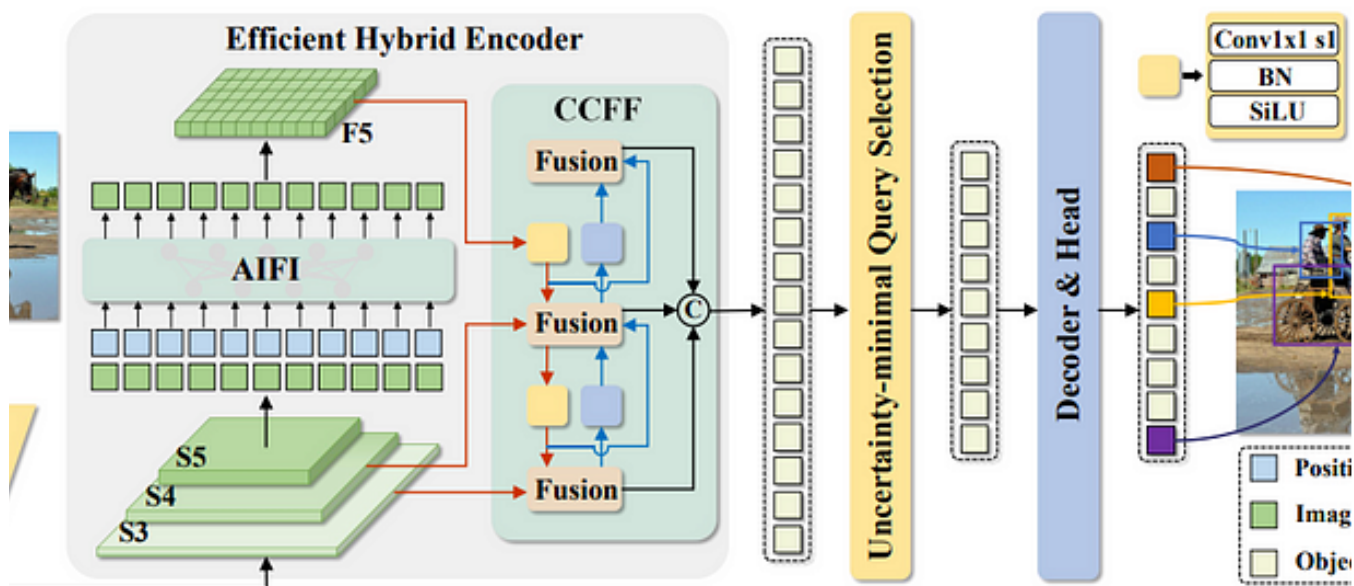
SegNet is a convolutional neural network (CNN) architecture designed for semantic segmentation in computer vision. Semantic segmentation is...

Oct 20, 2023 78 1



See all from Saba Hesaraki

## Recommended from Medium



Antonio Consiglio

## RT-DETR: A Faster Alternative to YOLO for Real-Time Object Detection (with Code)

Object detection has always faced a major challenge—balancing speed and accuracy. Traditional models like YOLO have been fast but...



Oct 28 172 1



**Stride**

The stride defines the step size with which the kernel moves across the input.

**Kernel**

A kernel, also known as a filter, is a small matrix of weights. The size of the kernel is usually 3x3

**Padding**

Padding refers to adding extra pixels around the input image, which helps preserve the spatial dimensions

**Max pooling**

Max pooling is a downsampling operation commonly used in Convolutional Neural Networks

In Convolutional Neural Networks (CNNs), a **convolution** is a fundamental



Sunny Kumar

## Building Blocks of CNNs: Comprehensive Layer Analysis with Formulas and Output Dimensions

A Convolutional Neural Network (CNN) is a type of deep learning model commonly used for analyzing visual data. Inspired by the way the...



Nov 10



### Lists



#### Natural Language Processing

1842 stories · 1465 saves



#### Practical Guides to Machine Learning

10 stories · 2068 saves



#### data science and AI

40 stories · 296 saves



#### Staff picks

778 stories · 1488 saves



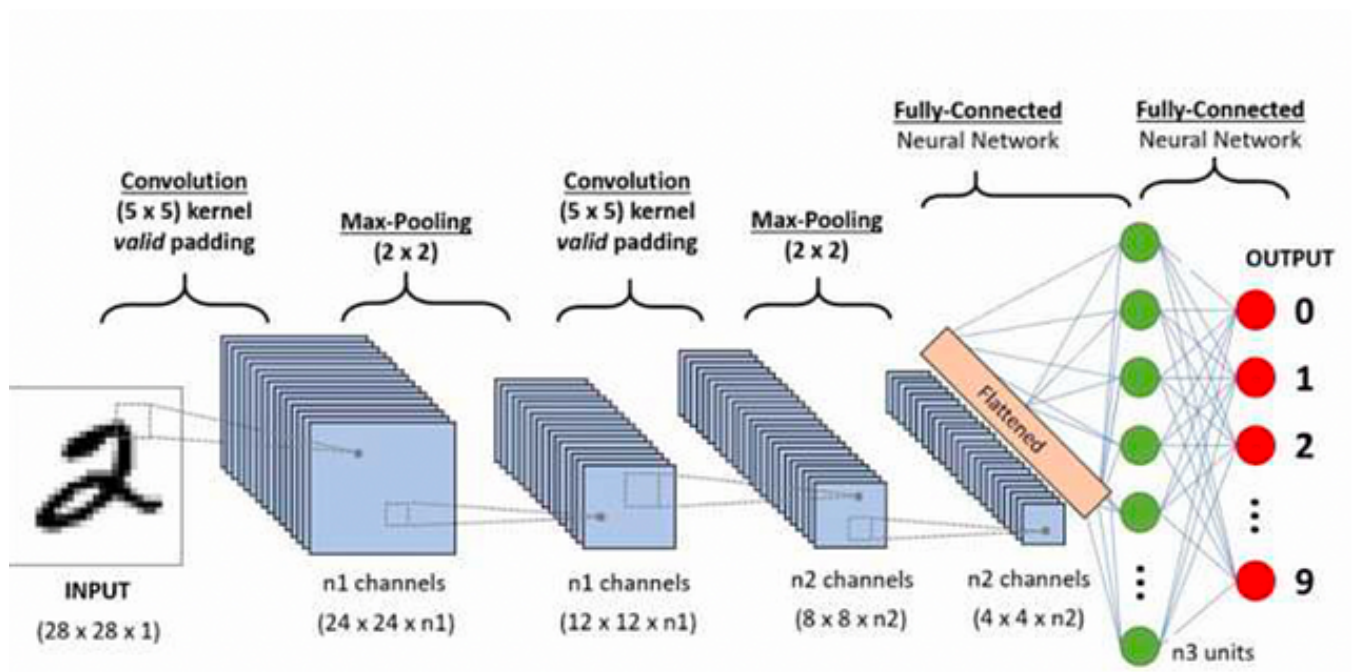


In The Deep Hub by Jorgecardete

## Convolutional Neural Networks: A Comprehensive Guide

Exploring the power of CNNs in image analysis

Feb 7 2.6K 38



Nivedita Bhadra

## Understand Convolution Neural Network (CNN)

Imagine you have a big picture and want to identify it—a cat, a dog, or a flower? A Convolution Neural Network(CNN) takes that picture...



Oct 30 117




## Notation

weights	$W_{\Lambda}[1] = 4 \times 3$	$W_{\Lambda}[2] = 2 \times 4$	$W_{\Lambda}[3] = 1 \times 2$	
layer sizes	$n_{\Lambda}[0] = 3$	$n_{\Lambda}[1] = 4$	$n_{\Lambda}[2] = 2$	$n_{\Lambda}[3] = 1$
biases	$b_{\Lambda}[1] = 4 \times 1$	$b_{\Lambda}[2] = 2 \times 1$	$b_{\Lambda}[3] = 1 \times 1$	
Activations	$A_{\Lambda}[1] = 4 \times m$	$A_{\Lambda}[2] = 2 \times m$	$A_{\Lambda}[3] = 1 \times m$	

$X = n \times m$  (transposed) matrix of input layer nodes across all training samples

$y_{\text{hat}} = 1 \times m$  row vector of predictions across all training examples

$A_{\Lambda}[l]$ : the  $n_{\Lambda}[l] \times m$  matrix that represents the values of activation outputs for nodes in a layer.




Aadil Mallick

## Learn to Build a Neural Network From Scratch—Yes, Really.

In this massive one hour tutorial, we're going to build a neural network from scratch and understand all the math along the way.

Sep 18 102 1



**tds** In Towards Data Science by Cristian Leo

## The Math Behind Convolutional Neural Networks

Dive into CNN, the backbone of Computer Vision, understand its mathematics, implement it from scratch, and explore its applications

★ Apr 9 🖱 873 💬 2



See more recommendations